Validation Feedback

The following sections present responses verbatim (lightly formatted for readability) to each survey item, attributed by participant ID. We conducted a member-checking study with original 31 participants, of which 11 responded (P_3, P_5, P_6, P_7, P_8, P_{12}, P_{16}, P_{22}, P_{23}, P_{25}, P_{30}). Participants received a copy of the anonymized paper and completed a feedback form covering accuracy of quote and overall analysis. Participants who wished to provide more detailed feedback could choose to (a) complete a 9-question in-form questionnaire, (b) participate in a follow-up interview, or (c) provide open-ended written responses. One (P_3) completed the in-form questionnaire; one (P_5) provided open-ended written responses; nine declined additional feedback beyond the core questions. In addition, 2 participants confirmed the accuracy over emails, but we do not include it here since they did not answer the survey questions.

Quote Accuracy and Concerns

Question: "If we quoted you in the paper, do you feel your quotes are accurately used in context?"

Responses:

- P_3, P_5, P_8, P_{12}, P_{22}, P_{23}, P_{30}: Accurately used in context.
- P_6, P_7: Mostly accurate with minor issues.
- P_{16}, P_{25}: Quotes not used in paper.

Question: "Please specify which quote and your concern for the quotes." (Conditional; shown to participants who reported issues.)

P_6: "It's not an accuracy issue, really, just a note. For the quote on page 7, about using an LLM to review the content from an LLM point of view…I never went beyond a couple days of testing the idea. I don't want to make it sound like this is something we consistently do! (Though I did get the idea from WTD where they may be people doing it on a regular basis.)"

P_7: "The passage begins with 'Less experienced participants reported relying on external style…' Then follows my quote. The text implies P7 may be in the 'less experienced' group, since the quote is used to illustrate that category. I am not looking for a change on this one, I am only reporting."

P_{22}: "No concern because the quotes were accurate."

Overall Analysis Accuracy

Question: "Overall, do you feel our analysis accurately represents your experiences and the practices you described in your interview?"

Responses:

- P_3, P_6, P_7, P_8, P_{12}, P_{16}, P_{22}, P_{23}, P_{25}, P_{30}: Yes, the analysis represents my experiences.
- P_5: Partially accurate with major issues.

Final Thoughts and Open-Ended Feedback

Question: "Please share any final thoughts, reactions, or feedback you have about the results in the paper." (Conditional; shown to participants who opted for open-ended response.)

P_5: "I take some issue with the assertion that 'documentation quality is achieved through a collaborative review process.' I think that the quality of documentation is more closely related to the writer's understanding of the functionality—so if the writer has explored the functionality independently (usually done before writing or during the drafting process), that is what is serving as the foundation for the content. If the writer cannot do so or has insufficient time / available resources to do so, the documentation clarity and quality will suffer directly.

I personally see review as a stopgap, but ultimately the trust lies with the writer and their expertise that they've built up effectively as an independent subject matter expert. This approach and standpoint might be more common in software documentation, and those that are technical enough to be practicing in 'docs like code' workflows but I can't say for certain.

In this way, the reviews for technical documentation function much like those of code reviews, and are often treated similarly—junior writers will see their work reviewed more thoroughly than that of a trusted partner of a senior writer, and so too will some engineers spend a lot more time closely reviewing content than others (as mentioned later on in the paper).

I also find it odd that your section on 'Presentation Quality' (starting line 484) doesn't mention style linters like Vale that can be used to automatically/programmatically/deterministically review for specific aspects of presentation quality, since the descriptions tend to mention primarily style content vs design content. I suppose it is mentioned later. Quality of Presentation Design tends to be evaluated visually by reviewing a preview of the page as it will look to customers and assessing skimmability and adherence to 'writing for the web' best practices.

The quotes from P3 and P21 throughout echo my experiences.

'API references, which are not systematically reviewed since they are directly generated from code, and therefore only way to discover issues is through complaints from users.' (lines 644–645) At my company and at others that I've worked at in the past, the API reference is reviewed and the spec is auto-generated from the code but there are docstrings added to the API spec and those get reviewed, often by a

writer. So findability can be an issue for API reference but that's more due to tool limitations I'd guess (like if you're building on Swagger directly vs interactive pages included in your own docs site, the way that tools like ReadMe or GitBook or Mintlify or Fern allow you to do).

Lines 752 you start talking about LLMs—I can say now that I use LLMs especially when constructing formulaic content, like repeated tasks in multiple examples, or when making updates to mark content as deprecated or in beta where I use the same sorts of reusable content references repeatedly and the LLM saves me from having to copy paste (I'm using the coding assistant Augment). We're planning to update the context.md of augment with style guidance to provide more correct prompts as we write, as well as using Vale for deterministic checks at PR review.

Lines 763, 'reported building the website locally using Sphinx or Docusaurus'—those tools don't help you build the website locally, but rather those are the tools that you use to build the documentation site generally, and those ones happen to let you also build it locally. Others don't let you do that but you can usually still stage the content in some way."

Review Stages

Question: "Do the review process described in the paper match your experience with documentation review practices?"

- P_3: Yes, closely matches.
- P_5: Partially matches.

Question: "Which stage descriptions accurately match your experience? (check all that apply)"

- P_3: Self Review; Technical Review; Editorial Review; Post-Publication Feedback.
- P_5: Self Review; Technical Review; Editorial Review; Play Testing; Post-Publication Feedback.

Question: "Please share any additional thoughts, reactions, or feedback about the review stages. What did we get right/wrong?"

P_3**:**

What you got right.

Self Review. "This is spot-on. The distinction between experienced writers relying on internalised knowledge versus newer writers consulting style guides resonates. The testing against actual software is also crucial—I run through my own procedures."

Technical Review. "Accurate, especially the tension you identified between engineers wanting to document every feature versus writers focusing on user workflows. The 'negotiation' aspect is real and ongoing. The observation about technical review serving as a safeguard against editorial changes removing important details is particularly insightful."

Editorial Review. "The 'singular organisational voice' concept is exactly right. The three UX quality concerns (broken links, narrative gaps, structure assumptions) are the bread and butter of what editorial reviewers catch."

Post-Publication Feedback. "This captures the reality well, including how difficult it is to get quality-focused feedback from users and how traditional feedback channels break down with LLM consumption."

Play Testing. "You're right that not everyone uses this stage formally. In my experience, when play testing happens, it's often: informal and opportunistic—grabbing a colleague who hasn't seen the content; resource-dependent—only possible when QA or customer success teams have bandwidth; product-type dependent—more common for complex tutorials or API documentation than simple feature docs. The 'curse of knowledge' insight rings true though."

What could be refined.

"Stage flexibility. While you mention stages can happen in different orders, I'd emphasise that they're often simultaneous and overlapping. For example, I might do self-review while incorporating technical feedback or editorial and technical reviews might happen in the same PR conversation."

"Reviewer expertise gaps. The finding about senior writers knowing more than their reviewers is crucial and perhaps underexplored. This creates a paradox where the people most qualified to review are the ones who wrote the content."

"Review fatigue. One challenge not fully captured is how reviewers become less engaged with longer documents or frequent review requests. The 'LGTM without reading' problem you mentioned is significant and real."

Missing elements.

"Version control integration. The social dynamics of review comments in GitHub/GitLab PRs create unique pressures and behaviours that shape how feedback is given and received."

"Review ownership. Who's responsible when reviews are incomplete or superficial? This ownership question affects quality outcomes significantly."

P_5**:** "I thought the play testing was somewhat broad—usually happens only for tutorials and it's pretty rare to get that deep of a review. I think timing of review is also important and maybe under-mentioned; making sure that the reviews and content drafting happens close to when the project was developed so you can get a solid technical review before the engineer has moved on to other projects. If you lag too far behind then you will struggle to get more time with the engineer, but if you collaborate with them as the project is developed then it is easier to build that relationship and leads to fewer issues getting reviews (like mentioned later on with coping strategies) too."

Tools Taxonomy

Question: "We mapped tools across two dimensions: when they are used (Proactive / Corrective / Reactive) and who can use them (Individual / Collaborative). Does this categorization make sense for how you think about tools?"

- P_3: Partially; some aspects make sense but others don't.
- P_5: No, this doesn't match how I think about tools.

Question: "Are there important tools or categories of tools we missed?"

P_3**:** "I'd say that some tool categories seem underrepresented or missing:

Analytics and measurement tools. Heatmap tools (Hotjar, Crazy Egg) for understanding user behaviour; search query analytics to identify content gaps; user journey tracking across documentation sites; A/B testing platforms for documentation variations.

Content management and workflow tools. Content planning tools (Notion, Jira) for managing documentation roadmaps; translation and localization platforms (Crowdin, Lokalise)—absent despite being critical for global products; screenshots and media tools (Snagit, Loom, CloudApp) for visual documentation. The absence of CMS discussion might reflect the study's participant base, which was drawn from technical writing communities that tend to favor docs-as-code approaches over traditional CMS solutions.

Communication and coordination tools. Slack/Teams integrations for real-time review discussions; documentation request systems for managing incoming work; status dashboards showing review progress across teams (Jira, Wrike).

Specialised domain tools. API testing tools (Insomnia, Postman) mentioned briefly but understated; compliance and audit tools for regulated industries; brand/legal review platforms for marketing-adjacent content (Frontify, etc.)."

P_5**:** "Not that I'm aware of—LLM reviewers like CodeRabbit are starting to be more popular, as well as coding agents like Augment, Cline, or Cursor being more popularly reused for assistance writing documentation (vs the chatbots like ChatGPT or Claude) but I think that's largely due to timing."

Question: "Please share any additional thoughts, reactions, or feedback about the tools used in the review process. What did we get right/wrong?"

P_3**:** "As mentioned, analytics and operational tools are largely absent, including user behaviour analytics (heatmaps, search queries, journey tracking), content management systems for planning and workflow coordination, and specialised domain tools like translation platforms for global products. The paper also misses cross-functional integration tools that connect documentation to other business functions—CRM integrations that link support tickets to documentation gaps, customer feedback aggregation systems, and performance monitoring tools that track which docs actually reduce support burden or drive conversions.

Real-time collaboration and coordination tools are understated, particularly Slack/Teams integrations for review discussions, documentation request management systems, and visual media creation tools (screenshots, video capture) that are essential for keeping documentation current. The research focuses heavily on writing and review tools while missing the broader operational ecosystem that documentation teams need to function effectively at scale, especially tools that help bridge the gap between technical writers and other stakeholders who may not be comfortable with developer-focused workflows like GitHub PRs."

P_5**:** "I think about tools largely functionally—what objective are they serving and where do I want to use them—as part of the drafting process (perhaps as a vscode extension) or as part of the review process (perhaps as a github action on a PR) or as a periodic process that is run occasionally (like a broken link checker)."

Challenges and Coping Strategies

Question: "Which challenges resonate most with your experience? (check all that apply)"

- P_3: Poor planning of documentation cycles (documentation lags behind development); Lack of dedicated documentation tooling (must use developer-focused tools); Tool maintenance burden (responsibility falls on individuals); LLM reliability issues (LLMs validate incomplete content); Competing priorities (reviewers prioritize other work over documentation review).
- P_5: Tool maintenance burden (responsibility falls on individuals); Competing priorities (reviewers prioritize other work over documentation review); Poor planning of documentation cycles (documentation lags behind development).

Question: "Do you use these coping strategies? (check all that apply)"

- P_3: Strategic communication (distributed reviews, targeted questions, real-time interaction); Institutional and social pressure (escalation, public accountability).
- P_5: Strategic communication (distributed reviews, targeted questions, real-time interaction); Institutional and social pressure (escalation, public accountability).

Question: "Please share any additional thoughts, reactions, or feedback about the challenges and coping strategies. What did we get right/wrong?"

P_3**:** "I think that the challenges and coping strategies are accurate and capture the daily reality of documentation work. The process challenges—poor planning cycles where documentation lags behind development, competing priorities that make reviewers see writers as

creating additional burden, and expertise gaps where senior writers often know more than their reviewers—all resonate strongly with my own real-world experience. The tooling challenges, particularly the lack of dedicated documentation tools forcing teams to use developer-focused systems like GitHub PRs with non-technical stakeholders, and the maintenance burden of custom solutions creating single points of failure, are good observations that drive many organisational tool decisions.

The coping strategies of strategic communication (distributed reviews, targeted questions, real-time interaction) and institutional/social pressure (escalation, public accountability) represent workarounds that technical writers have developed to navigate organisational dynamics. What the research captures well is the emotional labor and political navigation skills that documentation quality requires—skills not typically part of job descriptions but essential for success. However, the research could expand on the psychological toll of being 'the bad guy for saying this needs more attention' and how these challenges create conditions for burnout. The validation of these frustrations as systemic issues rather than individual failures is valuable and provides a framework for understanding why quality remains challenging despite technical writers' best efforts."

P_5**:** "I think that the writing process is perhaps more collaborative vs combative than it seemed in the article's framing. It is easy to get frustrated with folks not providing review to your content but everything has give and take, so it's also important to call out all the ways that documentation writers often have agency in negotiating available timelines ('this is my list of priorities, where does your project sit?' or 'I can't work on that right now due to xyz but I can fit it in next week, does that work?') and requesting product managers to commit to documentation priorities as part of planning processes—effectively identifying 'will I hold the release or will I go without docs to hit the deadline' and letting docs plan accordingly. Docs also makes judgment calls about documentation detail and fidelity to align with timelines, balancing customer success with availability of content so that they can have something, and then we can go back and add an example or something or a detailed tutorial later on when we have more time or resources to do so."